

COMMERCE ORDER FULFILMENT DATA PLATFORM

 End-to-End Data Engineering Pipeline using Python, SQLite, Apache Airflow and Power BI 

Presented By:
SHASHIKANTH

Post Graduate Program in Data Science
(Data Engineering)

Hero Vired



AGENDA



Business Problem



Objectives



Dataset



Technology Stack



ETL Architecture

Bronz layer

Silver layer

Gold layer



Airflow

SQLite



Dashboard

Insights



Conclusion

BUSINESS PROBLEM

 Multiple CSV Files

 Manual Reporting

 Order Tracking Challenges

 Delivery Monitoring

 Customer Satisfaction Analysis

Goal: Build an automated ETL pipeline to transform raw e-commerce data into analytics-ready insights.

OBJECTIVES



Data Ingestion



Data Cleaning



Gold Layer
Creation



Airflow
Automation



SQLite Storage



Power BI
Dashboard

Dataset Overview

Source: Olist Brazilian E-Commerce Dataset

 Orders
99,441

 Customers
99,441

 Products
32,951

 Payments
103,886

 Reviews
99,224

 Sellers
3,095

Tables Used:

- Orders
- Customers
- Products
- Payments
- Reviews
- Sellers
- Order Items

TECHNOLOGY STACK

Tools and Technologies Used Throughout the ETL Pipeline

 Python - ETL Development



 Pandas - Data Cleaning & Transformation



 SQLite - Data Storage



 Apache Airflow - Workflow Orchestration




 Power BI - Dashboard & Analytics



 GitHub - Version Control

Project Architecture



BRONZE LAYER

Raw Data Ingestion & Storage Layer



Raw CSV Files
(Orders, Customers, Products,
Payments, Reviews, Items, Seller)



Bronze Layer



Purpose



Store Raw Data
Ingestion



Preserve Original Data



No Data Modification



Key Activities



Raw CSV



Historical Storage



Source Integrity

SILVER LAYER

(Data Cleaning, Transformation & Validation)

📦 Bronze Layer → 🧹 Cleaning Process → ❤️ Silver Layer

🎯 Purpose

- 🧹 Clean Data
- ✓ Improve Quality
- 📊 Standardize Data
- 🛡️ Validate Records


⚙️ Key Activities

- ✗ Remove Duplicates
- 🔄 Convert Timestamps
- 🚫 Handle Missing Values
- ✅ Data Validation

💡 Clean, validated and standardized datasets are prepared for creating Gold Layer analytical tables.

GOLD LAYER

(Analytics-Ready Data Warehouse)



Purpose

-  Business Analytics
-  KPI Reporting
-  Performance Analysis
-  Dashboard Creation



Output

-  Star Schema
-  SQLite Database Tables
-  Power BI Ready
-  Optimized for Fast Queries

DATA QUALITY

(Ensuring Data Accuracy, Consistency, and Reliability)

Validation Checks

Validation Rule

Duplicate Order IDs
Missing Values Handling
Timestamp Format Validation
Negative Payment Values
Negative Freight Values
Data Type Standardization



Result


Passed
Completed
Passed
Passed
Passed
Completed

Key Outcome

All validation checks were successfully completed, producing clean, consistent, and analytics-ready data for the Gold Layer and Power BI Dashboard.

STAR SCHEMA





APACHE AIRFLOW

(Automating the ETL Pipeline Workflow)

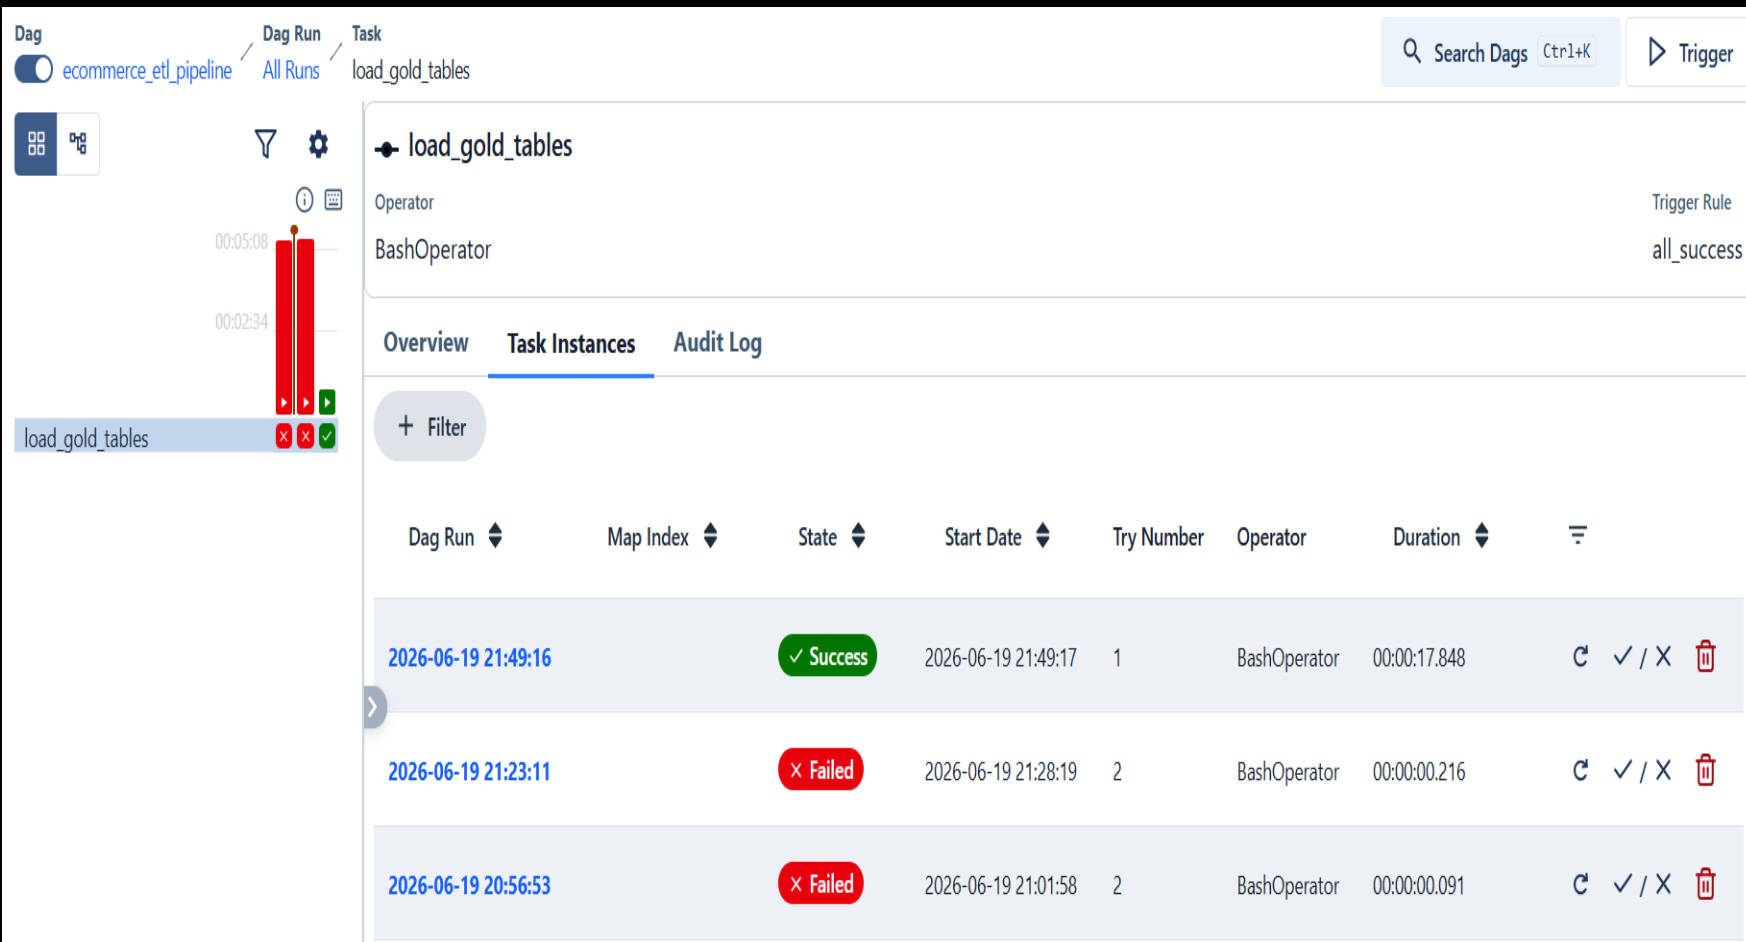


Key Features

- Workflow Automation
- DAG Scheduling
- Task Monitoring
- Logging & Retry



Airflow automates the execution of the ETL pipeline, ensuring reliable scheduling, monitoring, and logging.



DAG EXECUTION


ETL Pipeline Execution

Workflow Steps

- ① DAG Triggered
- ② Read Gold Layer CSV Files
- ③ Execute `load_to_sqlite.py`
- ④ Load Data into SQLite
- ⑤ Task Completed Successfully

The Airflow DAG successfully executed the ETL workflow by loading the Gold Layer datasets into the SQLite database with execution monitoring and task status tracking.

ETL PIPELINE SUCCESS




The screenshot shows the Apache Airflow interface for the DAG 'ecommerce_etl_pipeline'. The 'Runs' tab is active, displaying a table of DAG runs. The most recent run, 'manual_2026-06-19T16:19:16.298653+00:00', is in a 'Success' state. The table columns include Dag Run ID, Run After, State, Run Type, Triggering User, Start Date, Duration, and Dag Version(s).

Dag Run ID	Run After	State	Run Type	Triggering User	Start Date	Duration	Dag Version(s)
manual_2026-06-19T16:19:16.298653+00:00	2026-06-19 21:49:16	Success	Manual	admin	2026-06-19 21:49:17	00:00:09.312	v2

- DAG Executed Successfully
- Gold Layer Loaded
- SQLite Database Updated
- Dashboard Ready

-  ETL Pipeline Completed Successfully
-  Gold Layer CSV Files
- 
-  Apache Airflow DAG
- 
-  SQLite Database
- 
-  Power BI Dashboard

The ETL pipeline completed successfully, transforming analytics-ready Gold Layer datasets into the SQLite database for business intelligence reporting through Power BI.




SQLITE DATABASE (Centralized Storage for Gold Layer Tables)

-  **Database Features**
 - Lightweight Database
 - Stores Gold Layer Tables
 - Analytics Ready
 - Source for Power BI

💡 SQLite stores the transformed Gold Layer tables, providing a centralized data source for Power BI.

POWER BI DASHBOARD

(Interactive Business
Intelligence &
Performance Analytic



BUSINESS INSIGHTS

(Key Findings from the Power BI Dashboard)



Order Performance

99K+ Orders were successfully processed through the ETL pipeline, providing a comprehensive view of order fulfilment performance.



Revenue Analysis

The platform generated approximately **16.01 Million** in total revenue, indicating strong sales performance.



Delivery Performance

The average delivery time was **12.09 days**, reflecting an efficient order fulfilment process.



Customer Satisfaction

The average customer review score was **4.09/5**, showing a high level of customer satisfaction.



Payment Insights

Credit Card was identified as the most preferred payment method among customers.



Business Value

The interactive dashboard enables stakeholders to monitor KPIs, evaluate operational performance, and make data-driven business decisions.

CHALLENGES & SOLUTIONS

(Challenges Encountered During ETL Pipeline Development)

◆ Challenge 1

Multiple Raw Data Sources

Solution

Integrated multiple CSV datasets into a unified ETL pipeline using Python and Pandas.

◆ Challenge 2

Data Quality Issues

Solution

Performed duplicate removal, timestamp conversion, missing value handling, and validation checks during the Silver Layer transformation.

◆ Challenge 3

Workflow Automation

Solution

Implemented Apache Airflow to automate ETL execution, scheduling, monitoring, and logging.

◆ Challenge 4

Centralized Data Storage

Solution

Stored analytics-ready Gold Layer tables in SQLite, creating a centralized repository for Power BI reporting.

💡 Key Achievement

Successfully developed an end-to-end automated Data Engineering pipeline that transforms raw e-commerce data into business-ready insights.

FUTURE SCOPE

(Potential Enhancements for the ETL Pipeline)

Cloud Deployment

Deploy the ETL pipeline on cloud platforms such as Microsoft Azure, AWS, or Google Cloud Platform for improved scalability.

Big Data Processing

Integrate Apache Spark to process large-scale datasets with better performance and distributed computing.

Real-Time Data Processing

Implement real-time data ingestion and streaming using Apache Kafka or Apache Spark Streaming.

Predictive Analytics

Apply Machine Learning models to forecast sales, predict delivery delays, and analyze customer behavior.

Incremental Data Loading

Replace full data loads with incremental ETL processing to improve execution efficiency and reduce processing time.

Long-Term Vision

Develop a fully automated, cloud-based, real-time Data Engineering platform capable of supporting enterprise-scale analytics.

CONCLUSION

KEY ACHIEVEMENTS

- ✓ End-to-End ETL Pipeline
- ✓ Medallion Architecture
- ✓ Apache Airflow Automation
- ✓ SQLite Data Storage
- ✓ Interactive Power BI Dashboard

 Final Takeaway

"Raw Data → Business Insights"

THANK YOU

